

Divided CPU and Node Aware Elastic Net

L3P3

Thursday, November 13, 2014

This document summarizes the first approach in order to expand the node aware elastic net regularized predictor v3 to include resource awareness. This second iteration adds two columns to each input matrix consisting on the mean cpu value for the predicted node for each training instance time and the variation of that resource over the mean. The whole model creation is described in the file “rnen2-whole.R”, found in the private L3P3 Github repository. The models are exported in a variable called “models_sc”, which is a list of all the combinations of nodes and events in the system, containing each one the best created model for each event.

```
require(ggplot2)
```

```
## Loading required package: ggplot2
```

```
load("data.Rdata")
```

Our objective is now to study the performance of this models. To do so, we’re going to create a matrix where the rows are the system nodes and the columns are the events that happened on the system. The coefficients of the matrix will be the f-score of each model in each node.

Let’s now extract some insightful analytics.

- First, let’s see how many models were correctly created and how many models failed to create:

```
## [1] "Possible models:" "1335"
```

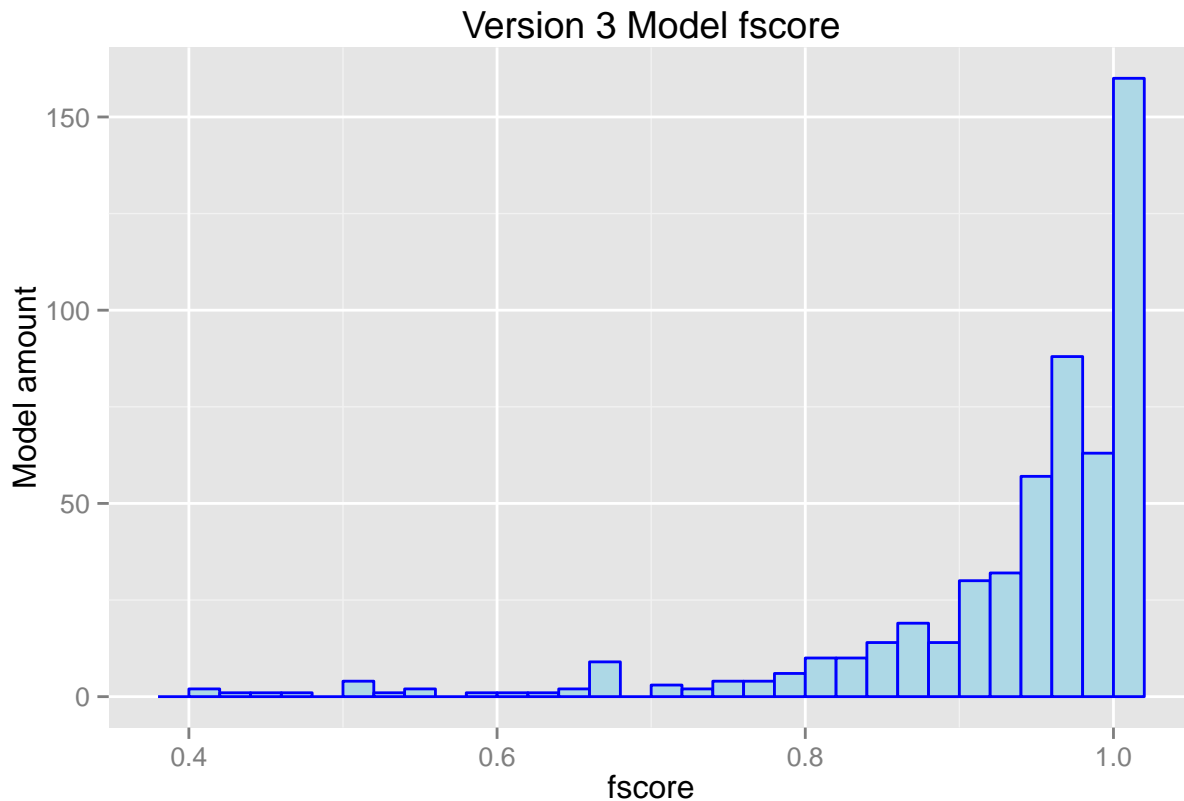
```
## [1] "Created models:" "542"
```

```
## [1] "Failed models:" "793"
```

The percentage of created models is a 40.5992509%. A logical result, as each model has now far fewer data to train.

- Now, we’ll study obtained fscores in the 563 correctly created models:

```
## stat_bin: binwidth defaulted to range/30. Use 'binwidth = x' to adjust this.
```



The histogram shows good results for the models that actually work. A 79.3357934% (430 models) is at least 0.9 and a 91.697417% (497 models) is at least 0.8. The average obtained fscore of the created models is 0.9318644, a slightly lower amount compared to the one found in the previous iteration. These models are hugely simpler than those, though, which makes still a feasible option in deployment.

- Lastly, we will study which critical events were correctly modelled and their fscore. There are three different critical events on the system, which we will study separately:
- 96731154: this event indicates a critical threshold violation. It only appears in two nodes.

```
## fwpeuib02.sceu.corp fwpeuid02.sceu.corp
##                      0                      1
```

- 68917: this event's message points that a device has stopped responding to external requests and/or polls. It is greatly associated with a change in the network on the 1st of June.

```
##      aswpeuib02.sceu.corp      aswpeuib04.sceu.corp
##                      0.00                      0.00
##      aswpeuid01.sceu.corp      aswpeuid02.sceu.corp
##                      0.00                      0.00
##      aswpeuid03.sceu.corp      aswpeuid04.sceu.corp
##                      0.00                      1.00
##      aswpeuin01.sceu.corp      aswpeuin02.sceu.corp
##                      0.00                      0.96
##      aswpeuin03.sceu.corp      aswpeuin06.sceu.corp
##                      0.00                      0.00
```

##	aswpeuin08.sceu.corp	aswpeuin09.sceu.corp
##	0.00	0.00
##	aswpeuin10.sceu.corp	aswpeuin12.sceu.corp
##	0.00	0.00
##	ASWPEUIN13.sceu.corp	ASWPEUIN19.sceu.corp
##	0.00	0.00
##	dnsrespeuin02-adm.sceu.corp	dswpeuin01.sceu.corp
##	0.00	0.00
##	dswpeuin02.sceu.corp	dswpeuin03.sceu.corp
##	1.00	0.00
##	dswpeuin04.sceu.corp	dswpeuwt01.sceu.corp
##	0.00	1.00
##	dswpeuwt02.sceu.corp	dswpeuwt03.sceu.corp
##	0.00	0.00
##	dswpeuwt04.sceu.corp	fwpeuid02.sceu.corp
##	0.00	0.00

A total of 26 models are created, out of which 4 were correctly created, this is, a 15.3846154%, with an average fscore of 0.99 and a standard deviation of 0.01.

* 69481:this event's message poitns that a whole chassis has stopped responding to polls. While it is g

##	aswpeuib02.sceu.corp	aswpeuib04.sceu.corp	aswpeuid01.sceu.corp
##	0.0000000	0.0000000	0.0000000
##	aswpeuid02.sceu.corp	aswpeuid03.sceu.corp	aswpeuid04.sceu.corp
##	1.0000000	0.0000000	0.0000000
##	aswpeuin01.sceu.corp	aswpeuin02.sceu.corp	aswpeuin03.sceu.corp
##	0.0000000	0.0000000	0.0000000
##	aswpeuin06.sceu.corp	aswpeuin08.sceu.corp	aswpeuin10.sceu.corp
##	1.0000000	0.0000000	0.0000000
##	aswpeuin12.sceu.corp	ASWPEUIN13.sceu.corp	ASWPEUIN19.sceu.corp
##	0.0000000	0.0000000	1.0000000
##	dswpeuin01.sceu.corp	dswpeuin02.sceu.corp	dswpeuin03.sceu.corp
##	1.0000000	0.9615385	0.0000000
##	dswpeuin04.sceu.corp	dswpeuwt01.sceu.corp	dswpeuwt02.sceu.corp
##	1.0000000	0.9787234	0.9787234
##	dswpeuwt03.sceu.corp	dswpeuwt04.sceu.corp	
##	0.0000000	0.0000000	

A total of 23 models are created, out of which 8 were correctly created, this is, a 34.7826087%, with an average fscore of 0.9898732 and a standard deviation of 0.0074744.

Conclusions

Adding the CPU consumption has increased the amount of created models and also increased (though in a really small amount) the average fscore, but has decreased the performance of the critical events models. In a model with 1400 input variables it is expected for a single variable not to cause a huge change in performance. It is probable to see a larger change with every resource consumption variable.